

AI Agent Plan

High-Level Roadmap

May 2026

Presented by Matt Stone, CTO

WHAT WE WANT TO ACHIEVE

AI agents across the customer lifecycle – bringing customers in, serving them, and running the business – plus a master engineering agent that builds and operates the platform itself. The aim is a scalable, always-on, lower-cost EPM experience, with a licensed human adviser retained where it matters and engineering velocity multiplied.

Akane	Customer acquisition – web chat, eligibility, structured intake → handoff to a human adviser.	First
Misato	Customer communications & service – status, statements and support for active customers.	Later
Rie	Back-office & operations – document verification, settlement, reconciliation.	Internal
Yumi	Investment account – setup, monitoring, end-of-life settlement.	Internal
Motoko	Engineering & operations – the master agent. Builds and runs the platform and the four product agents above. Human-supervised.	Live today

Codenames are internal only – customers always see "the FutureProof assistant". Motoko is a different class from the four product agents: it builds and operates them, and is not part of the customer journey.

THE APPROACH

- **Build on what we have** – the agents sit on the existing calculation engine and EPM model, not a new platform.
- **Start where value is provable** – the customer-acquisition agent (Akane) and the engineering agent (Motoko, already delivering); expand to communications, operations and investment as each proves out.
- **Rent intelligence, own the sensitive tier** – a frontier cloud model for the customer-facing surface; a local model on our own hardware for sensitive and internal data.
- **Keep a human in the loop** – a licensed adviser confirms before any advice is issued.

HOW WE MITIGATE RISKS

- **Stage-gated** – we commit one stage at a time, each with go and kill criteria; investment grows only as evidence supports it.
- **Data & compliance** – sensitive data stays onshore / in-house; customer data flows only through compliant, no-training channels with privacy-counsel sign-off; the human-adviser gate keeps us inside the advice rules.
- **Technology & dependency** – the master agent is human-supervised (no autonomous production changes); owning local hardware hedges reliance on any single cloud-AI provider.
- **No over-promising** – the broader AI-ecosystem position is treated as upside to earn, not a moat we claim today.

ROUGH BUDGET

Deliberately modest to start: roughly one engineer's time for a quarter, a small monthly cloud-AI bill, and one Mac Studio (~A\$8k, 96GB) – fundable within the current round. Investment scales only as each stage's gate is passed. Detailed costings, success metrics, and the build-vs-buy specifics sit in the companion plan.