



AI Agent Plan

What we're building, how we'll measure it, and what it costs

AUDIENCE

Management

PRESENTED BY

Matt Stone, CTO

STATUS

Plan

DATE

May 2026

THE PLAN, IN ONE SLIDE

We are building AI agents on our existing model – staged, gated, and measurable.

- **Five agents** – four across the customer lifecycle, plus a master engineering agent (Motoko) that builds and runs the platform and is **already delivering today**.
- **Stage 1 first** – the customer-acquisition chat (Akane), an API so other AIs can call our EPM calculator, and the engineering agent. We commit Stage 1 only; the rest is earned at hard gates.
- **Build vs buy** – rent frontier intelligence for the customer-facing surface; own one Mac Studio for sensitive and internal data.
- **Why now** – customers increasingly start financial decisions with an AI assistant, and being callable by those assistants is becoming table stakes. We act on that, but commit only to what is provable.

Real deliverables, real costs, real metrics, and a clear build-vs-buy decision – not an ecosystem land-grab we cannot yet evidence.

THE AGENTS

Five agents: four run the business, one builds it.

Akane	Customer acquisition – web chat (EPM FAQs, eligibility), structured intake → compliance gate → handoff to a licensed human adviser.	Build first
Misato	Customer communications & service – status, statements and support for active customers.	Later stage
Rie	Back-office & operations – document verification, settlement, reconciliation, internal queries.	Partial · internal
Yumi	Investment account – setup, monitoring, end-of-life settlement.	Partial · internal
Motoko	Engineering & operations – the master agent. Builds and runs the platform and the four product agents above. Human-supervised.	Live today

Codenames are internal only – customers always see "the FutureProof assistant". Motoko is a different class from the four product agents: it builds and operates them, and is not part of the customer journey.

WHERE WE ARE TODAY

Honestly: one agent live, two partial, two to build.

Working today

- The EPM calculation engine and actuarial model
- Chat infrastructure in the app (channels, message store)
- **Motoko** – the engineering/ops agent is already delivering on this codebase

Partial / internal & not yet built

- **Rie** and **Yumi** – internal tooling exists, not productionised
- **Akane** (customer chat) – to build (Stage 1)
- **Misato** (customer comms) – later stage
- The agent-to-agent API (EPM schema) – to build (Stage 1)

No vapourware: the slide above is the real state of the codebase, not a roadmap of intentions.

STAGE 1 – WHAT WE BUILD

Three concrete deliverables, on what we already have.

1 · Akane Layer 1 – customer chat

Web chat on the existing calculation engine: EPM FAQs, eligibility, and a structured intake → compliance gate → handoff to a licensed human adviser. Runs on a frontier model in an AU region, identifiers minimised.

2 · The EPM API (Layer 2)

Publish an MCP + OpenAPI spec for the EPM calculator so other AI assistants can find and call it. Modest, concrete, and the foundation for any future ecosystem position.

3 · The engineering agent (Motoko)

Continue and formalise the engineering/ops agent already delivering – the most proven, measurable AI leverage we have today.

Not in Stage 1

Misato (customer comms), the L3 ecosystem partnerships, and any autonomous production changes. These come later, behind gates – or are treated as options to earn.

HOW WE'LL KNOW IT WORKED

Stage 2 proceeds **only if all of these hold** – otherwise we de-scope or kill.

≥ 40%

intake completion (of engaged visitors)

≥ 4/5

adviser-rated handoff quality

cost per qualified lead

< current channel**< 5s**

p50 customer response time

zero

compliance breaches in pilot

engineering cycle-time

down ≥ 25%

These are the Stage-1 targets; exact thresholds are calibrated against current-channel baselines before the pilot. Every gate has a kill criterion as well as a go criterion.

Outsource frontier intelligence, or run large models ourselves?

Outsource – frontier API (Claude / GPT / Gemini)

- **Quality/safety:** frontier-class – best for a customer advice-adjacent surface
- **Cost:** ~\$0.30-1.50 per conversation → ~\$1-2k/mo at launch volume
- **Latency:** 3-6s responses
- **Data:** leaves the building – mitigable (AU region, zero-retention, identifier-minimisation)
- **Ops:** none; live in weeks

Own – Mac Studio + open model

- **Hardware:** 96GB (~A\$8k) runs 70B-class; 512GB (~A\$15k) only for the largest open models
- **Quality:** below the latest frontier for hard reasoning & guardrails
- **Latency:** ~15-20s for a full reply; a handful of concurrent sessions per machine
- **Data:** never leaves the building – strongest compliance story
- **Ops:** we run the stack, uptime, updates

INFRASTRUCTURE – THE DECISION

Rent intelligence, own the sensitive tier.

Customer-facing → frontier API

AU region, zero-retention / no-training, identifiers minimised before the call. Best quality and safety; ~\$1-2k/mo at launch – far cheaper than self-hosting at this stage.

Sensitive / internal → one Mac Studio

A local 70B-class model for raw-PII workloads, internal documents, batch and fine-tuning. Data never leaves the building; doubles as the dev box and a compliance hedge.

Crossover (defined, not religion): revisit local-for-customer-facing only if sustained API spend exceeds ~\$8-10k/month ($\approx$ tens of thousands of conversations/month) **and** an open model passes our evaluation bar. Reviewed quarterly.

DATA & LEGAL POSITION

A data-classification policy – not "trust the vendor".

- **Accountability stays with us.** Under Privacy Act APP 8 we remain responsible for any offshore handling – sending data to a third party does not offload liability.
- **Defensible default:** AU-region processing + enterprise zero-retention / no-training terms + **identifier-minimisation before the call.**
- **Strictest classes stay in-house** – raw PII and sensitive financial records route to the local Mac Studio. A classification policy decides what goes where.
- **Advice stays compliant** – the human-adviser gate means the agent never issues personal advice on its own.
- **Sign-off before go-live** – privacy counsel, and APRA CPS 230 (material service provider) checks if applicable.

Landscape, not legal advice – the specific obligations (APP 8 / s16C, ACL/AFSL, CPS 230, the 2024 Privacy reforms on automated decisions) are confirmed with counsel.

THE STAGE PLAN

We commit Stage 1; the rest is earned.

Stage 1 Months 0–6 · committed	Akane Layer 1 + the EPM API (Layer 2) + the engineering agent. <i>Gate</i> → <i>Stage 2</i> : the Stage-1 metrics hold (intake, handoff quality, cost/lead, latency, compliance, eng cycle-time).
Stage 2 Months 6–12	Akane to general availability; Misato (customer comms); first partner conversations. <i>Gate</i> → <i>Stage 3</i> : signed partner LOIs + agent-attributed conversion tracking to plan. <i>De-scope</i> to Akane-only if partner appetite is weak.
Stage 3 Year 2–3	Scale (UK/US) and pursue the ecosystem position (an option, not a promise). <i>Continue only</i> on material agent-attributed acquisition.

Quarterly AI-landscape review. Programme commitment grows only as evidence supports it.

WHAT IT COSTS

Modest to start; it scales only as gates pass.

~1

engineer-quarter
(Stage 1 build)

~\$1-2k

per month
(frontier API, launch)

~A\$8k

one-off
(one Mac Studio, 96GB)

Stage 1 is fundable within the current round. Stages 2-3 are re-costed and re-approved at each gate; we do not pre-commit budget beyond the next stage. Detailed costings sit in the build plan.

RISKS – AND HOW WE CONTAIN THEM

Each risk has a concrete mitigation already in the plan.

AI quality / safety on the customer surface

Frontier model + bounded scope + the human-adviser gate before any advice. Pilot before GA.

Data & compliance

Classification policy: AU-region/enterprise terms for de-identified; in-house for sensitive. Counsel + CPS 230 sign-off.

Vendor / dependency

The local Mac Studio is a working hedge; the architecture stays portable across model providers.

Over-investment

Stage gates with explicit kill criteria; commitment never runs ahead of evidence.

THE HONEST CUT

What's real now – and what's an option to earn.

Real now – shipping value

- The engineering/ops agent (Motoko) – delivering today
- Akane Layer 1 – the Stage-1 build, on infrastructure we control
- The EPM API – a concrete, modest deliverable

An option to earn – not a moat we claim

- Adoption of the EPM schema by other AI assistants
- Partnerships with adviser AI and super / pension-fund AI
- We pursue this and track it – we do not pitch it as a moat we already hold

WHAT WE'RE DOING

This is the plan we're executing.

- **Building Stage 1** – Akane Layer 1, the EPM API, and the engineering agent – time-boxed and gated on the metrics shown.
- **Adopting the hybrid infrastructure** – frontier API for customer-facing, one Mac Studio for the sensitive/internal tier – with a data-classification policy and counsel sign-off.
- **Buying one Mac Studio** (~A\$8k, 96GB) for the sensitive-internal tier and development.
- **Treating the ecosystem position as an option** – pursued and tracked quarterly, not claimed as a moat.

| *Commit to what's provable now; earn the rest.*